

RECURRENT NEURAL NETWORKS

Vincent Barra

LIMOS, UMR 6158 CNRS, Université Clermont Auvergne

SEQUENTIAL DATA

Sequential data

- ▶ Timeseries data (temperature, pressure, stock market...)
- ▶ Speech / music
- ▶ Videos
- ▶ ...



Problem!

- ▶ Arbitrary length
- ▶ Huge number of parameter for a model ?

PROPERTIES

Need for memory

- ▶ Data in a sequence is not identically, independently distributed
- ▶ Need for a context, thus for memory

Paul is a small boy and is hungry. He goes to the restaurant and eats a lot.

PROPERTIES

Need for memory

- ▶ Data in a sequence is not identically, independently distributed
- ▶ Need for a context, thus for memory


Paul is a small boy and is hungry. He goes to the restaurant and eats a lot.

THE question

How to model sequential data, context and memory ?


DATA REPRESENTATION

Neural nets need numerical values, not words




DATA REPRESENTATION

Neural nets need numerical values, not words



DATA REPRESENTATION

Neural nets need numerical values, not words




Corpus

Paul, small, hungry,
he, goes, restau-
rant, and, lot,...

DATA REPRESENTATION

Neural nets need numerical values, not words



Corpus


Paul, small, hungry,
he, goes, restau-
rant, and, lot,...

Indexing

a → 1
boy → 2
...
small → 16

DATA REPRESENTATION

Neural nets need numerical values, not words



Corpus

Paul, small, hungry,
he, goes, restau-
rant, and, lot,...

Indexing

a \rightarrow 1
boy \rightarrow 2
...
small \rightarrow 16

One-hot encoding

a = (1, 0, ..., 0)
boy = (0, 1, 0, ..., 0)
...
small = (0, 0, ..., 1)

DATA REPRESENTATION

Why not directly using the index as a descriptor ?

Example : distance between "a" and "small"

- 1 Indexes: $d^2("a", "small") = (16 - 1)^2 = 225$
- 2 One hot encoding: $d^2("a", "small")^2 = 2$

- ▶ Indexes : distance depends on the values of the index
- ▶ One hot encoding : whatever two different words, they have the same distance if they are different

EMBEDDINGS

Word2vec

Learns word embeddings by estimating the likelihood that a given word is surrounded by other words.
Bag of words, skip Gram.


Dimensions

- ▶ One-hot vectors:
high-dimensional and
sparse
- ▶ word embeddings:
low-dimensional and
dense.

Generalization


- ▶ One-hot vectors:
constrained by the
corpus
- ▶ word embeddings:
Generalization,
capabilities.

PROCESSING INDIVIDUAL DATA POINT




"Paul"


PROCESSING INDIVIDUAL DATA POINT




PROCESSING INDIVIDUAL DATA POINT



PROCESSING INDIVIDUAL DATA POINT




PROCESSING INDIVIDUAL DATA POINT



$$\hat{y}_t = f(x_t)$$

PROCESSING INDIVIDUAL DATA POINT




$$\hat{y}_t = f(x_t)$$

t	Paul	is	a	small	...	He	goes	to	the	...
	0	1	2	3	...	8	9	10	11	...

x_8 depends on $x_0 \Rightarrow$ with this model, no possible relation.


INTUITION: NEURONS WITH RECURRENCE




$$\hat{y}_t = f(x_t, h_{t-1})$$

- ▶ x_t : input
- ▶ \hat{y}_t : output
- ▶ h_{t-1} : past memory

FOLDED VERSION



RECURRENT NEURAL NETWORKS




Apply a recurrence relation each time step to process a sequence

$$(\forall t \geq 0) \quad h_t = f_W(x_t, h_{t-1})$$

- ▶ h_t : current cell state
- ▶ f_W : neural network with parameter matrix W
- ▶ x_t : input
- ▶ h_{t-1} : old cell state (memory)

To keep memory, W is shared through time.

RECURRENT NEURAL NETWORKS



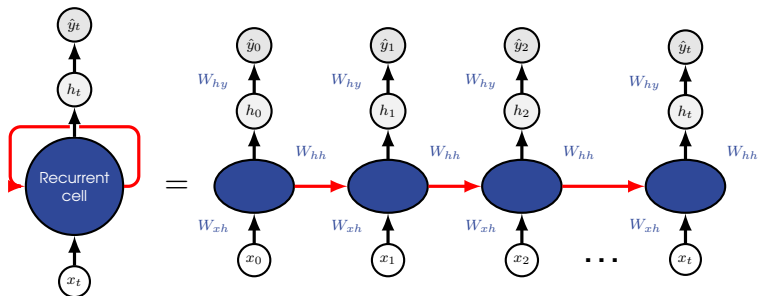
- 1 Update hidden state

$$h_t = \tanh(W_{xh}^T x_t + W_{hh}^T h_{t-1})$$


- 2 Compute output vector

$$\hat{y}_t = W_{hy}^T h_t$$

FOLDED VERSION- FORWARD PASS



FOLDED VERSION- FORWARD PASS



KERAS IMPLEMENTATION FROM SCRATCH

```
class RNNFromScratch(tf.keras.layers.Layer):  
    def __init__(self, nb_units, dim_in, dim_out):  
        super(RNNFromScratch, self).__init__()  
        self.Whh = self.add_weights([nb_units, nb_units])  
        self.Wxh = self.add_weights([nb_units, dim_in])  
        self.Why = self.add_weights([dim_out, nb_units])  
  
        self.h = tf.zeros([nb_units, 1])  
  
    def call(self, x):  
        self.h = tf.math.tanh(self.Wxh*x + self.Whh*self.h)  
        y = self.Why*self.h  
  
        return y, self.h
```



KERAS IMPLEMENTATION: SIMPLERNN

SimpleRNN layer


SimpleRNN class

```
tf.keras.layers.SimpleRNN(  
    units,  
    activation="tanh",  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    recurrent_initializer="orthogonal",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    recurrent_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    recurrent_constraint=None,  
    bias_constraint=None,  
    dropout=0.0,  
    recurrent_dropout=0.0,  
    return_sequences=False,  
    return_state=False,  
    go_backwards=False,  
    stateful=False,  
    unroll=False,  
    **kwargs  
)
```


SOME ARCHITECTURES




One to many
Captioning, music generation



Many to one
Sentiment analysis




Many to Many
Video annotation



Many to many
Translation

UNFOLDED VERSION- BACKWARD PASS



UNFOLDED VERSION- BACKWARD PASS

Computing the gradient w.r.t. h_0 involves

- ▶ many factors of W_{hh}
- ▶ repeated gradient computation

Many high values

- ▶ Exploding gradients
- ▶ Gradient clipping
- ⇒ Bouncing and unstable optimization

In all cases, possibility to loose long-term dependencies.

UNFOLDED VERSION- BACKWARD PASS

Computing the gradient w.r.t. h_0 involves

- ▶ many factors of \mathbf{W}_{hh}
- ▶ repeated gradient computation

Many high values

- ▶ Exploding gradients
- ▶ Gradient clipping
- ⇒ Bouncing and unstable optimization

Many small values

- ▶ Vanishing gradients
- ⇒ No gradient at all

In all cases, possibility to loose long-term dependencies.

BACKPROPAGATION THROUGH TIME

BPTT

- ▶ Basically chain rule as in classical backpropagation
- ▶ a bit more tricky, since gradients survive over time


Implementation

Already implemented, in Keras, using the classical `train` method.

WHAT'S NEXT ?

Limits

- 1
 - ▶ x_1 goes t times through \tanh .
 - ⇒ The influence of x_1 is much lesser than the one of x_t
 - ▶ W is shared: the weight associated to h_1 doesn't compensate this loss of memory
 - ⇒ RNN → short memory (the deeper the shorter)
- 2 Vanishing/exploding gradients



Limits

Some alternatives, improvements: LSTM, GRU... See next lecture !